

RedeployMate — Lean Canvas (B2B + B2C) & Moat/IP Options

Companion to the RedeployMate Project Brief. Updated 7 Sept with web-search verification — see inline “Verified” callouts for what was checked against current sources, and the note on “Job Care” for the one thing that couldn't be confirmed.

Part 1 — Lean Canvas: Customer, Problem, Solution, Value Proposition

Two tracks, because a B2B buyer and a B2C end-user need genuinely different problem framing and value props even though the underlying product (Evidence Agent + optional Reflective Companion) is the same. This matches Kevin's own suggestion to run B2B as the primary bet and B2C as an awareness play.

Track A — B2B (primary, per the pivot)

Customer	Priority	Problem	Solution (packaging)	Value Proposition
Turas Nua	● High — real lead in motion	Caseworkers assess clients but have no visibility once someone leaves the room until the next appointment; a caseworker has already asked for a one-pager.	Caseworker-facing tool used in or between sessions; caseworker keeps the one-page output on file. Their job board and appointments stay theirs.	“We don't touch what you already run — we make the weeks between appointments count, and give you something to show the DSP for it.”
Seetec	● High — verified equivalent, natural next call	Identical structure to Turas Nua (same state contract, Dublin/Border/West region) — same caseworker time-pressure problem.	Same tool and packaging as Turas Nua; reuse the one-pager with region-specific framing.	Same value proposition as Turas Nua.
Outplacement firms	● Medium — already identified, not yet actioned	Paid on placement outcomes but have thin per-candidate time; hard to prove impact between paid sessions.	White-label Evidence Agent + cheat-board embedded in the firm's existing coaching workflow.	Better placement/completion metrics to report to corporate clients, without added headcount.
Universities (careers services)	● Medium	Need measurable graduate/alumni outcomes for funding and rankings; can't give every student 1:1 time.	Self-serve tool via the careers portal, with an advisor dashboard for flagged/at-risk students.	Demonstrable outcome data for funders and rankings, at near-zero marginal cost per student.
Recruitment agencies	● Medium	Spend billable time re-explaining a candidate's experience to multiple clients.	Recruiter-facing version turning a CV into a reusable two-role/evidence output.	Faster placement, more consistent candidate story across submissions.
Corporate HR & L&D / internal mobility	● High — new, promising	Considering internal redeployment before redundancy, but existing tools (Gloat, Fuel50, TechWolf) need a pre-built job architecture most mid-market teams don't have.	Lightweight, no-architecture-needed version for a single restructuring event, not an enterprise platform rollout.	Redeploy staff internally in weeks, not the 1–6 month onboarding competitors require.

Customer	Priority	Problem	Solution (packaging)	Value Proposition
Trade unions (SIPTU, Fórsa)	● High — new, promising, verified precedent	Represent members facing redundancy; want a concrete offer without becoming a training provider themselves.	Member-facing tool distributed through existing union training channels (e.g. SIPTU College).	A tangible, modern member benefit that costs the union nothing to deliver at scale.
EAP / employee benefits providers	● Exploratory	Already sell layoff-support benefits but want differentiated content without building it themselves.	Distribution partnership — RedeployMate becomes a module inside an existing benefit.	A more concrete, AI-native offering to differentiate their package.
Regional Skills Fora / LEOs / Intreo case officers	● Medium — real but slower cycle	Same caseworker-time problem as Turas Nua/Seetec, via the direct state channel.	Same caseworker tool, positioned as complementary to existing appointments.	Same “gap-filling, not replacing” pitch, at state scale.
Redundancy/income-protection insurers	● Exploratory	Some policies include “return to work” riders they must fulfil but don't want to build themselves.	Co-funding or white-label arrangement, similar to the EAP model.	Ready-made way to fulfil an existing policy benefit without new headcount.
Professional membership bodies	● Exploratory	Members facing mid-career transition; body wants to be seen supporting members through change.	Member portal integration, similar to the union model.	Membership-retention value: “we're with you even if your employer isn't.”
Independent career coaches / solo consultants	● Exploratory	Want the credibility of a “tool” for their practice without building one; smaller ticket than a firm contract.	Individual practitioner licence/subscription.	A concrete, AI-backed artifact to hand clients — differentiates a solo practice.
University alumni offices	● Exploratory	Alumni hitting redundancy years post-graduation are underserved, high-goodwill, tied to fundraising.	Alumni-facing version via alumni relations, not the student careers office.	Reinforces “we support you for life” — a fundraising/engagement asset, not just careers.
Retail/hospitality chains (rolling redundancies)	● Exploratory	Automation/store closures create repeated, larger redundancy events; HR shared-services need scalable support.	Procurement-led contract, larger volume, standard HR shared-services integration.	Consistent, scalable support without expanding the HR team each time.
International equivalents (UK/EU activation providers)	● Exploratory — later-stage	Same structural problem as Turas Nua/Seetec, in other markets.	Same product, localised course/taxonomy data per market.	Same value proposition, proven first in Ireland before expansion.

Shared thread across every B2B row: the handoff is where value leaks — an intermediary assesses someone, hands over a list, and has no visibility or leverage until the next appointment. RedeployMate fills that gap; it never replaces the intermediary's own relationship, board, or appointments.

Track B — B2C (secondary, “awareness play” per Kevin)

Customer	Priority	Problem	Solution (packaging)	Value Proposition
Redundant / restructuring-affected workers	● High — original design target	Assessed, handed a course list, then ~3 silent weeks until the next appointment — the stall.	Free self-serve Evidence Agent: CV/experience in, two reachable roles + one step this week out.	“You already did the work. We just help you prove it — for free, without waiting three weeks.”
People on mandatory jobseeker programmes	● High — same core segment, via Turas Nua/Seetec path	Same gap problem, compounded by mandatory-engagement pressure.	Same tool, positioned as something to bring <i>into</i> the mandatory appointment.	Walks into the next appointment already holding something concrete.
Over-50s workers	● Medium	Most experience, least legible “evidence” in keyword-matched systems; face age-bias.	Same Evidence Agent, tuned to surface transferable skills from long tenures.	Turns decades of experience into provable, current evidence rather than a liability.
Women returners	● Medium	Career-break gap reads as a red flag to ATS/recruiters.	Same tool, confirms the work already done — “no extra story about the break.”	Confirms real capability without needing to explain or apologise for the gap.
Voluntary career changers	● Medium	Proactively anticipating industry decline or burnout — a much larger pool than reactive redundancy.	Same free tool, marketed outside formal redundancy channels.	Get ahead of the stall before it happens.
Recent graduates (AI displacing entry-level roles)	● Medium — Kevin flagged as a bigger, adjacent problem	Increasingly can't get the first “real” evidence-building job at all.	Reachable via university careers services; helps translate coursework/projects into evidence.	Turns coursework into evidence when there's no work history yet.
Immigrants / relocators	● High — sharpest B2C addition	Overseas experience doesn't map cleanly to Irish-recognised roles.	Same tool, tuned for cross-border qualification/experience translation.	Makes real overseas experience legible to Irish employers without re-credentialing from scratch.
Gig / freelance workers	● Exploratory	Want to formalise into employment but have non-traditional “evidence” (client work, not job titles).	Same tool, adapted to accept project/client-based evidence.	Converts freelance/gig work into evidence a traditional employer recognises.
Workers in sunset roles	● Exploratory	Anticipate disruption (e.g. automation) before redundancy actually happens.	Same free tool, positioned as early-warning/preparation, not crisis response.	Get the reachable-roles picture before the redundancy letter arrives.

Shared thread across every B2C row: the emotional core (per Andrew's original draft) is the stall — plans that don't survive the gap between appointments. The B2C product is positioned as a funnel/awareness asset per Kevin's suggestion, not the primary revenue line.

Quick synthesis: why these priorities

The  High rows above aren't arbitrary: **Turas Nua** is real and in motion; **Seetec** is its verified structural twin, so it's the natural next call once Turas Nua is validated; **trade unions** and **corporate internal-mobility** are the two “new” segments with a clear budget holder and plausible near-term urgency; and **immigrants/relocators** is the sharpest B2C addition because it's arguably a purer version of the core problem than “redundant worker” in general.

Verified stats worth folding into the pitch (confirmed via search, 7 Sept):

- **FIT** (Fastrack into Information Technology): over 250,000 jobseekers have completed FIT programmes to date; 120,000 have secured quality employment. Founded 1999.
- **Women ReBOOT** (Technology Ireland Digital Skillnet): over 500 women re-entered the workforce; post-12-week placement, over 90% of participants have secured full-time roles at companies including ESB, AIB, Dell, Guidewire, and Bank of Ireland.
- **Generation Ireland**: 200–250 people supported into work annually via 10–12 bootcamps/year, 1,000+ total since 2019; participants are paired with a mentor **through the programme and up to 6 months after** — this is a direct, verified match for the “gap after the programme” channel insight from the 3 Sept mentor call.
- **On “Job Care”** (named in the 3 Sept notes as a comparator alongside Generation Ireland and Technology Ireland): this didn't turn up in search under that name — it may be a mishearing/mis-transcription in the auto-generated notes. Worth asking Brigitte/whoever attended to confirm the actual name before it goes in front of anyone external.
- **On the €50bn/9% GDP stat** already in the pitch: confirmed as a real PwC figure, but it's from PwC's **2023** Women in Work Index (based on 2021 participation data). PwC's newer **March 2026** Women in Work Index shows Ireland climbing to 11th of 33 OECD countries (the biggest improvement in the OECD, driven by a narrowing gender pay gap), but also flags that **female full-time employment has fallen and female unemployment has risen** across the OECD this year — arguably an even stronger, more current pitch line (“this is getting worse, not solved”) than the 2023 GDP figure alone.

Process & Design Decisions — Aligned with the Team Presentation

Added 7 Sept so this doc and the slide deck shared with the group stay in sync. Status per item reflects where things stood after the working session — treat as our proposed position going into group discussion, not a closed decision.

The process: today vs. with RedeployMate

The comparison now forks at step one, not after it — a change from the original design (which treated the intake as unchanged and shared between both paths):

- **Today (unchanged, status quo)**: Assessed (basic caseworker check-in) → Handed a list (course list or role suggestions) → ~3 silent weeks (no contact until the next appointment) → often: the plan doesn't survive the gap.
- **With RedeployMate (redesigned)**: Full Assessment (CV in, 1–2 questions, more if the CV is thin) → Two Real Roles (not a job board — roles they can reach) → Walk Out With Proof (a plan, a course, likely interview questions) → they already have something concrete.

Decision: the intake becomes a full assessment

Our proposal, going to the group for confirmation: the CV comes in, and RedeployMate asks 1–2 short questions — more only if the CV is thin or unclear. This reuses the same flex-by-thinness logic as the interview-question design decision below, just applied at intake instead of a separate later step. It moves value earlier: richer evidence exists before the caseworker even sits down.

- **Before — passive assessment:** caseworker reviews the CV as-is; person leaves with a list; no adaptive follow-up.
- **Now — full assessment:** CV in, 1–2 adaptive questions (more if thin), confirmed by the person — feeding the caseworker's own session, not replacing it.
- **Worth watching:** keep this as RedeployMate's own intake feeding the caseworker's session, not overwriting it — otherwise it risks quietly undercutting the “we don't replace what you run” positioning that the whole buyer-gain case rests on. As long as it stays a pre-step that hands the caseworker richer material rather than a substitute for their meeting, the positioning holds.

Two open design decisions — status

Both of these were flagged as proposals for the group; internally, both are agreed:

- **How many interview questions?** Not a fixed number for everyone — ask only about what the CV genuinely doesn't evidence. **Agreed:** flex the question count by how thin the CV already is. Two modes: a quick 2–3-question pass for a live caseworker session, an optional deeper pass for someone using the free version alone at home.
- **How many roles to show?** “Two” is deliberate — it's the line against being a job board, keeps quality/vetting high, fits a 90-second demo. Real limitation: a hard cap of two can hide the actual best-fit role and makes bias harder to audit. **Agreed:** compute and rank the top 5–10 candidate roles internally (for bias auditing and fallback), keep the default person-facing view at two, with a “show me more” option for segments that genuinely want breadth.

The Unique Value Proposition — the ABC through-line

“Everyone else hands you a list. We hand you proof — two roles you can already reach, confirmed by you, with a board to walk into the interview holding. And we work for the caseworkers and unions who'll never have a Gloat-style platform to plug into, because we don't need one.”

Three things make this genuinely hard to copy, not just well-phrased — each contrasted against a specific competitor category already mapped above:

1. **The output is proof, not a list.** Turas Nua, FIT, ReBOOT, Generation Ireland all end with training or a course list. None produce a personalised, evidence-confirmed “two reachable roles + gap-closing plan + interview cheat-board.” Copying this means building the whole evidence-to-interview pipeline, not adding a chatbot to an existing course catalogue.
2. **It works where the enterprise tools structurally can't, without months of setup.** Gloat, Fuel50, and TechWolf all require an existing job architecture/skills taxonomy before they add value — a verified “sequencing problem” for anyone without one. Activation providers, unions, and outplacement firms don't have that infrastructure and won't build it. RedeployMate runs per-candidate against an external taxonomy (ESCO), so it's usable on day one for exactly the buyers who need it — a gap in the market's structure, not just a feature gap.

3. **Trust is designed in, not bolted on.** Nothing is used until the person confirms every line — the actual reason a caseworker would let this near their client relationship, since it can't quietly override their judgement. In a market where “the algorithm decided” is a real adoption blocker, that's a genuine edge, not an ethics footnote.

Honest caveat: none of this is unbreachable on day one — a well-resourced competitor could in principle build all three. What can't be built quickly is real outcome data (which confirmed-evidence → gap-plan → course sequences actually led to a placement) from a live partner like Turas Nua or Seetec. That only exists once the relationship is running for a while — it's what turns “better designed” into “actually more accurate over time,” and it's the honest version of a moat at this stage, not a claim the tech itself is unclonable.

Evidence for the Pitch — Why “Evidence, Not Keywords” Holds Up

Supporting research for the claim that people (and the CV format itself) routinely fail to capture someone's real capability — the core justification for the Evidence Agent's skill-surfacing mechanism. Confirmed via search, 7 Sept. Confidence level noted per item; use accordingly.

△ Flag first — a stat to actively avoid: the widely-repeated claim that “men apply for a job at 60% qualified, women only at 100%” (from Sheryl Sandberg's *Lean In*, repeated in HBR and Forbes) traces back to a **speculative comment by one Hewlett-Packard executive, not real data**. When the Behavioural Insights Team tested it properly, the actual gap was small (56% vs. a slightly higher figure for men). A 2024 peer-reviewed study (Salwender et al., *European Journal of Social Psychology*, 4 studies) found **no significant gender difference** in one experiment and only a weak, non-robust one in others. This is exactly the kind of “sounds true, isn't” stat Andrew's own rule was written to catch — if it's anywhere in the current pitch pack, pull it.

What actually holds up, from strongest to more caveated:

Finding	Source	Confidence
Despite 80% of returners actively upskilling before job-hunting, 99% still find re-entry a genuine struggle; two-thirds of professional female returners accept roles well below their actual level — driven by reduced confidence leading people to undervalue their own skills and self-select into junior roles.	Career Returners Indicator 2026 (annual survey, 12 years / 5,000+ returners supported)	Strongest — current (2026), large sample, directly matches the women-returners persona
Recruiters were barely better than a coin flip (55% accuracy) at predicting who'd pass a technical interview based on resume review alone, and specifically underestimated the weakest-looking candidates by an average of 23 percentage points — meaning real ability was often meaningfully higher than the resume suggested.	interviewing.io analysis of real recruiter/candidate outcome data	Strong — real outcome data, though a single platform's dataset, not peer-reviewed
Women under-report soft skills (leadership, time-management) and omit achievements from resumes more than men do.	Adzuna study, reported by HR Dive	Solid — a named, real study via a reputable trade outlet
Self-assessed skill level vs. actual test performance: 56% of users underestimated their own skill level, 32% overestimated, only 11% were accurate.	Workera platform analysis	Real and quantified, but vendor-published (Workera is a skills-assessment company) — not independent academic research

How to use this in the pitch: lead with the Career Returners Indicator 2026 figure for the women-returners persona specifically (it's the most current and most direct), use the interviewing.io finding as the general “CVs are a bad instrument, regardless of who's underselling” argument, and treat the Workera and Adzuna figures as supporting colour rather than headline stats given their sourcing caveats. Do not use the 60%/100% stat under any framing.

Part 2 — Moat / IP: Pros & Cons of the Three Options

The question on the table: pursue “**agentic orchestration**” framing, pursue a **fine-tuned model**, or **neither** (a different defensibility story). Judged against what actually matters for a 2-week hackathon build with the team's current resources (small eng team of ~2, prompt-based credits only, no confirmed GCP/GPU access, ~6 days left to 13 Sept).

Option A — “Agentic orchestration” as the moat

Pros

- Achievable now — it's largely a framing and architecture exercise on top of what's already built (extract → confirm → question → ESCO-map → sequence → predict is already an orchestrated multi-step pipeline).
- Matches what Kevin explicitly suggested, so it's aligned with how the judges (Scale Ireland) are likely to read “AI-native.”
- Demoable live in 90 seconds — the pipeline itself is visible, which plays directly to the “Build with AI” judging criterion.
- **Verified:** LangGraph specifically ships human-in-the-loop as a core, built-in capability (not something you'd have to hand-roll), and is already used in production by Klarna, Uber, and J.P. Morgan — so “agentic orchestration with a human confirmation gate” is a technically accurate, credible framing, not just marketing language.
- The workflow design, guardrails, confirm-before-use gates, and ESCO-mapping heuristics are genuine (if soft) IP — harder to casually copy than “a course list,” even though the underlying LLM is commodity.

Cons

- Weak as a *real* moat — LangGraph and similar frameworks are open-source, well documented, and explicitly designed to be easy to adopt; a competent competitor, or another hackathon team hearing the same mentor advice, could replicate the pipeline shape in weeks.
- No data or network effects on its own — orchestration logic alone doesn't compound over time without outcome data feeding back into it.
- Risks reading as a buzzword reframe rather than substance if a judge asks the follow-up question directly: “what stops Turas Nua's own dev team building this next quarter?”

Option B — Fine-tuned model

Pros

- Genuinely harder to replicate — a model tuned on real Irish labour-market data (validated CV→evidence patterns, real outcomes) would be a proprietary asset, not just a prompt.

- Strong technical credibility and “wow factor” for the “Build with AI” score.
- Could plausibly improve real product quality over time and reduce long-run inference cost versus paying for frontier-model calls at scale — relevant to the commercial-viability criterion.
- **Correction after checking current costs:** LoRA fine-tuning of a small open-source model (e.g. via Together AI or Fireworks) is genuinely cheap in 2026 — commonly under \$500, sometimes under \$50, and doesn't require GCP/GPU infrastructure at all; several teams report going from a few thousand labelled examples to a working fine-tune for well under \$500. So the “we can't afford the compute” objection from the original version of this analysis was overstated — cost and infrastructure are no longer the real blocker.

Cons

- **The real blocker is data, not compute.** Fine-tuning needs a genuinely labelled training set — real CVs mapped to real evidence, real skill gaps, real course outcomes — and the team only has synthetic CVs and no validated outcome data yet. Assembling and quality-checking even a few hundred good examples in the days remaining is a real task on its own, separate from the (now cheap) training step.
- Even with cheap training, current guidance from practitioners increasingly favours **well-engineered prompts over fine-tuning** for exactly this kind of task — a frontier model with a good system prompt and RAG may already match or beat a rushed fine-tune trained on thin data.
- Risk of a “fake it” demo — claiming a fine-tuned model built on synthetic data only is thin, and judges may probe on data provenance.
- Directly collides with the Challenge 26 rubric's own responsible-AI warning — “do not infer protected traits or rank human worth... test for occupational and demographic bias” — a rushed fine-tune on unvalidated data is exactly where that risk shows up, and reviewers may ask about bias testing you won't have had time to do properly.
- Still adds new scope and a model/infra decision very late in the build, competing directly with finishing the actual app and the Turas Nua pitch.

Option C — Neither: trust/relationship & data moat

Pros

- Already built, honestly — the “confirm every line, nothing used until you say yes” workflow is a real, working feature, not a claim.
- Directly answers Kevin's actual objection from a different angle: the moat isn't the AI plumbing, it's (a) being the layer caseworkers trust enough to actually use because it doesn't threaten their role, and (b) the outcome/relationship data that accumulates once a partner like Turas Nua is live — a classic, credible early-stage B2B moat (“thin today, compounds through the relationship”).
- Turas Nua's own moat is largely a state contract and relationships, not superior technology — this framing meets a B2B judge where incumbents in this exact market actually compete.
- Zero new engineering scope this close to the deadline — time instead goes to the Turas Nua pitch, the Screen 1 fix, and rehearsing the story.

Cons

- Less “AI-native” sounding for a hackathon explicitly themed around agents — risks under-scoring “Build with AI” if presented on its own with nothing technical alongside it.
- Trust/relationship moats are real but slow to prove in a 90-second demo; harder to visualise than a pipeline diagram or a fine-tuned model claim.
- Doesn't fully close the “anyone could rebuild the course list” critique — it reframes the question rather than answering it head-on.

At a glance

Option	Feasible by 13 Sept?	Judge-legible as “AI depth”?	Real long-term defensibility?	Main risk
A. Agentic orchestration	Yes — mostly framing/cleanup	Yes	Weak alone	Reads as buzzword if not paired with something else
B. Fine-tuned model	Cost/infra no longer the blocker (LoRA fine-tunes now run under \$500) — but data is: no real labelled dataset exists yet	Yes, if real	Strong, if trained on real outcome data	High risk of a thin “fake it” demo on synthetic data; still competes for time against finishing the app
C. Trust/relationship + data moat	Yes — already exists	Weak alone	Genuinely strongest, if the Turas Nua relationship lands	Undersells AI depth if not paired with A

My read, for the team to weigh — not a final call

These three aren't mutually exclusive, and picking just one probably isn't the strongest answer. The combination that fits both the timeline and the actual objection Kevin raised:

- **Lead with A (agentic orchestration) as the technical story** — it's true, it's demoable, and it's what the room is primed to respond to. It's also technically accurate: the human-confirmation gate already built into the demo is a legitimate agentic-orchestration pattern (LangGraph itself ships this as a first-class feature), not a fig leaf.
- **Back it with C (trust workflow + the Turas Nua relationship) as the commercial defensibility story** — this is the honest answer to “why can't they just build this,” and it's already true rather than aspirational.
- **Park B (fine-tuning) as an explicit phase-2 roadmap line**, not something built now — and the reason has shifted since the first pass on this: it's not that fine-tuning is too expensive or needs infrastructure the team doesn't have (that's no longer true in 2026), it's that there's no real labelled dataset yet. The honest roadmap line is: “once we have real placement outcome data from a partner like Turas Nua, that's what we'd fine-tune on — and doing it then, cheaply, is realistic.” That directly ties the roadmap back to the data moat in C, so it reads as a coherent story rather than a bolt-on.

That said, this is a judgement call that depends on how the Turas Nua conversation actually goes — worth revisiting once you know whether that relationship is real traction or just a polite meeting.