

RedeployMate

Everyone else hands you a list. We hand you proof.

Working Session — Full Detail for the Build Team


TechIreland National AI Challenge 2026

What's in this deck

- The process — today, side by side with RedeployMate
- The evidence behind the interview step, and its real risks
- Two open design decisions we want your read on
- Who we sell to, the value case, and our technical proposal

The Process: Today vs. With RedeployMate

Redesigned per the group's steer: even the first step is different now — not just what happens after it.



Decision: Step 1 Becomes a Full Assessment

Resolved per the group's steer — bringing it to the session to confirm, not re-open.

OUR PROPOSAL

The CV comes in, and RedeployMate asks 1–2 short questions — more only if the CV is thin or unclear.

This reuses the exact same flex-by-thinness logic already agreed for the interview step (Slide 6) — applied right at intake instead of waiting for a separate later step. It moves value earlier: richer evidence exists before the caseworker even sits down.



Before — Passive Assessment

Caseworker reviews the CV as-is. Person leaves with a list. No adaptive follow-up.



Now — Full Assessment

CV in, 1–2 adaptive questions (more if thin), confirmed by the person — feeding the caseworker's own session, not replacing it.

Worth watching: keep this as RedeployMate's own intake feeding the caseworker's session, not overwriting it — so “we don't replace what you run” (Slide 10) still holds.

Why the Interview Step Matters — The Evidence

A CV alone routinely under-represents what someone can actually do. This isn't an opinion:

99%

of career returners still find re-entry a genuine struggle — even after actively upskilling.

Career Returners Indicator, 2026

55%

recruiter accuracy predicting who passes an interview, from a resume alone — barely better than a coin flip.

interviewing.io analysis of real hiring outcomes

More supporting evidence — weaker sourcing, still worth knowing

Women under-report soft skills (leadership, time-management) and omit achievements more than men do.

Adzuna study, via HR Dive

Solid — named study

Self-assessed skill vs. actual test result: 56% of users underestimated their own skill level, only 11% were accurate.

Workera platform analysis

Vendor-published, not academic

⚠ Don't use: the famous “men apply at 60% qualified, women only at 100%” stat is a myth — traced to a speculative HP comment, not data. Rigorously tested (BIT + a 2024 peer-reviewed study) and found not to hold up.

Would This Actually Work? Weighing It Honestly

The interview step is arguably the actual product — which means it deserves real scrutiny, not just enthusiasm.



PROS

- Matches real coaching practice — people chronically undersell what they've actually done
- Adaptive questioning scales effort to real need, not a one-size-fits-all form
- The confirm-every-line step is a genuine safety net — it catches bad inferences before they're used
- Surfaces exactly the skills the CV format itself is bad at capturing (see evidence slide)



CONS — real, not hypothetical

- Hallucination risk: an inferred skill that doesn't hold up sends someone into a real interview with false confidence — worse than a generic course list
- Question design can quietly bias who benefits — favouring articulate written English risks under-serving immigrants, over-50s, and returners specifically
- Tone risk: even 2–3 questions can feel like an interrogation to a stressed population if worded badly
- Honest gap: nobody has piloted whether AI-elicited “overlooked skills” actually hold up in real interviews yet — strong hypothesis, not proven

Two Decisions We Want Your Read On

Both framed as proposals — not settled, open for the group.



How Many Questions?

Not a fixed number for everyone.

Ask only about what the CV genuinely doesn't evidence — a rigid “10 questions for all” wastes time on strong CVs and under-probes thin ones.

PROPOSAL

Flex the question count by how thin the CV already is — thin evidence gets more questions, strong evidence gets fewer or none. Two modes: a quick 2–3-question pass for a live caseworker session, and an optional deeper pass for someone using the free version alone at home.

Question design matters more than question count — see the pros/cons slide before this. This logic is now built into the redesigned intake too (Slide 3).



How Many Roles to Show?

“Two” is deliberate, not arbitrary.

It's the line against being a job board, keeps quality/vetting high, and fits a 90-second demo. But a hard cap of two can hide the actual best-fit role, and makes bias harder to audit.

PROPOSAL

Compute and rank the top 5–10 candidate roles internally (for bias auditing and fallback), but keep the default person-facing view at two — with a simple “show me more” option for segments (careers offices, proactive changers) that genuinely want breadth.

Who It's For, and Why It Works

The B2B lens — our near-term focus, per the pivot away from B2C/government



CUSTOMER

Employment activation providers, outplacement firms, unions. Turas Nua & Seetec caseworkers are the near-term focus — a real conversation already in progress.



PROBLEM

A caseworker assesses someone, then has zero visibility until the next appointment — weeks away. That gap is where value leaks, for the client and for the caseworker's own outcomes.



SOLUTION

A caseworker-facing Evidence Agent that fills the gap: confirmed evidence → two reachable roles → interview-ready board. Their job board and relationship stay theirs.



VALUE PROPOSITION

“We don't replace what you already run — we make the weeks between appointments count, and give you proof to show for it.”

A free, self-serve B2C version runs in parallel as an awareness channel — not the primary revenue line. Full segment map: next slide.

The Full Segment Map — Every Buyer We Mapped

● *High priority* ● *Medium* ● *Exploratory*

B2B — 15 segments

- **Turas Nua** — *real lead in motion*
- **Seetec** — *verified twin, natural next call*
- **Outplacement firms** — *paid on outcomes, thin time/candidate*
- **Universities (careers)** — *need graduate outcome data*
- **Recruitment agencies** — *reusable candidate story*
- **Corporate HR / internal mobility** — *no architecture needed, unlike Gloat*
- **Trade unions (SIPTU, Fórsa)** — *verified training precedent*
- **EAP / benefits providers** — *distribution partner*
- **Regional Skills Fora / LEOs / Intreo** — *same problem, state channel*
- **Redundancy insurers** — *white-label “return to work” riders*
- **Professional bodies** — *membership retention*
- **Independent coaches** — *solo-practice licence*
- **University alumni offices** — *fundraising/engagement asset*
- **Retail/hospitality chains** — *procurement-led, larger scale*
- **International (UK/EU equivalents)** — *later-stage, ESCO already EU-wide*

B2C — 9 segments

- **Redundant/restructuring workers** — *original design target*
- **Mandatory jobseeker programme people** — *same path via Turas Nua/Seetec*
- **Over-50s workers** — *experience, least legible evidence*
- **Women returners** — *confirms work, no explaining the gap*
- **Voluntary career changers** — *bigger pool than reactive redundancy*
- **Recent graduates (AI-displaced)** — *Kevin's flagged adjacent problem*
- **Immigrants / relocators** — *sharpest B2C addition*
- **Gig / freelance workers** — *non-traditional evidence*
- **Workers in sunset roles** — *proactive, not reactive*

Detail on each segment's problem/solution/value framing lives in the companion Lean Canvas doc.

UNIQUE VALUE PROPOSITION

Everyone else hands you a list. We hand you proof.



Proof, Not a List

Evidence-confirmed roles plus an interview board
— not another course catalogue.



Zero Setup Needed

No existing job architecture required — unlike
Gloat, Fuel50, or TechWolf.



Trust, Designed In

Nothing is used until the person confirms it
themselves.

Where we make a difference: the part of this market too small and un-wired for the big HR platforms to bother with.

This is the product-level moat. The next slide translates it into what the buyer actually gains.

The Gain for the Businesses Who Buy This

What's actually in it for Turas Nua, Seetec, an outplacement firm, or a union — concretely.



Time Back for Caseworkers

Less time re-explaining the same CV across sessions — the evidence work is already done when they sit down.



Better Numbers to Report Upward

Completion and placement metrics improve because the gap weeks aren't wasted — something concrete to show funders or clients.



Zero Disruption to What They Run

Their job board, appointments, and relationship stay exactly as they are. This fills the gap — it doesn't touch their system.



A Defensible AI Story, Not a Liability

Nothing is used until the client confirms it. If a funder or regulator asks how the AI is controlled, there's a real answer.

Directional, not yet measured — no live partner data exists until a pilot runs. Framed as reasoned value drivers, not promised ROI.

How Do We Stay Hard to Copy? Three Options We Weighed

Kevin's flag: the course-list feature is thin and quickly copyable. Here's what we considered.

A. Agentic Orchestration

PRO

Feasible now — largely framing on what's already built. Demoable, matches judge expectations.

CON

Weak alone — LangGraph-style pipelines are open-source; a competitor (or another team) could replicate the shape in weeks.

B. Fine-Tuned Model

PRO

Strong defensibility if real — costs dropped a lot (~\$50–500 via LoRA, no GCP needed).

CON

Blocked by data, not cost — no real labelled outcomes exist yet; synthetic-only risks a thin “fake it” demo.

C. Trust + Data Moat

PRO


Genuinely strongest long-term — already built (confirm-before-use), compounds once a partner is live.

CON

Slow to prove, and undersells AI depth on its own — needs pairing with A to land with judges.

Agentic Orchestration Now. Fine-Tuning as Phase 2.

Lead with A (feasible, demoable) — back it with C (the Turas Nua relationship) — park B until real data exists. Bringing this to the group for agreement, not announcing it as decided.



NOW

- Agentic orchestration pipeline, LangGraph-style
- Human-in-the-loop confirmation gates on every inference
- ESCO-based role mapping
- Demoable and true today — not a roadmap slide



PHASE 2 — NOT BUILT YET

- Fine-tune once a partner (e.g. Turas Nua) is live and generating real outcome data
- Cheap once real data exists — roughly \$50–500 via LoRA, no GCP needed
- The blocker is real data, not cost or compute

Over to You

Points for the group to review and agree on

1

Step 1 is now a full assessment: CV in, 1–2 questions, more if thin. (Slides 2–3)

2

Flex the interview question count by how thin the CV is. (Slide 6)

3

Keep “two roles” shown, rank 5–10 internally for audit/fallback. (Slide 6)

4

Does the Unique Value Proposition and buyer-gain framing actually land? (Slides 9–10)

5

Comfortable committing to B2B-first (Turas Nua / Seetec) as the near-term focus? (Slide 7)

6

Agentic orchestration now, fine-tuning as Phase 2. (Slide 12)